Linear models and misspecification

Statistics III - CdL SSE

Tommaso Rigon
Università degli Studi di Milano-Bicocca

Homepage

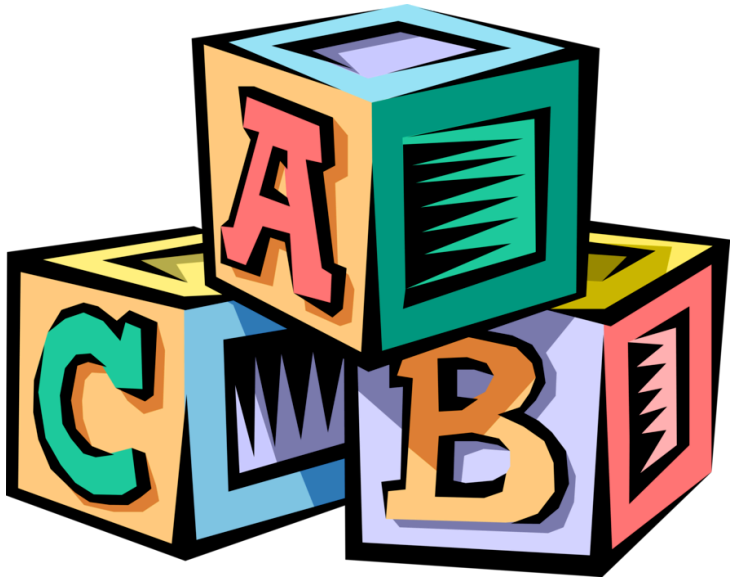

"Everything should be made as simple as possible, but not simpler"

Attributed to **Albert Einstein**

- This unit will cover the following **topics**:
  - **Recap**: linear models and the modeling process
  - Robustness of OLS estimates, sandwich estimators
  - Weighted least squares
  - Box-Cox transform, variance stabilizing transformations
- The main theme is: what should we do when the **assumptions** of linear models are **violated**?
- We will push the linear model to its limit, using it even when is not supposed to work.
- The symbol 📖 means that a few extra steps are discussed in the **handwritten notes**.

> The content of this Unit is covered in **Chapter 1** of Salvan et al. (2020). Alternatively, see **Chapter 2** of Agresti (2015) or **Chapter 5** of Azzalini (2008).

The modeling process

Car data (diesel or gas)

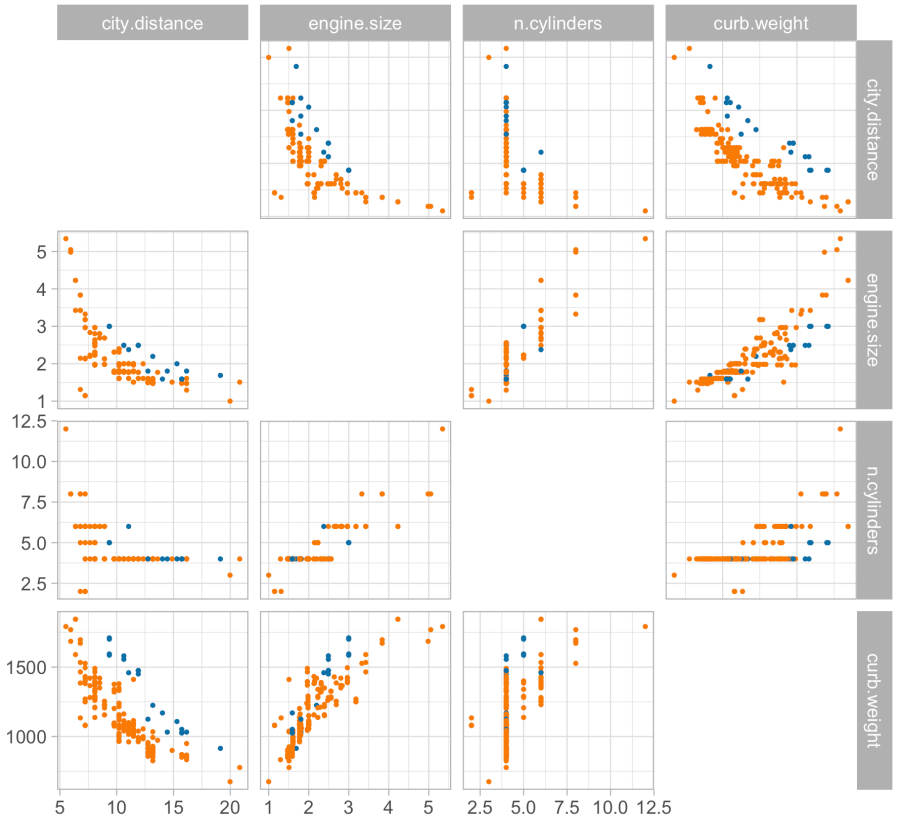

- We consider data for $n = 203$ models of cars in circulation in 1985 in the USA.
- We want to **predict** the distance per unit of fuel as a function of the vehicle features.
- We consider the following **variables**:
  - The city distance per unit of fuel (km/L, `city.distance`)
  - The engine size (L, `engine.size`)
  - The number of cylinders (`n.cylinders`)
  - The curb weight (kg, `curb.weight`)
  - The fuel type (gasoline or diesel, `fuel`).

We assume you are already familiar with linear models. The following is a brief recap rather than a full discussion.

Home page

Linear regression

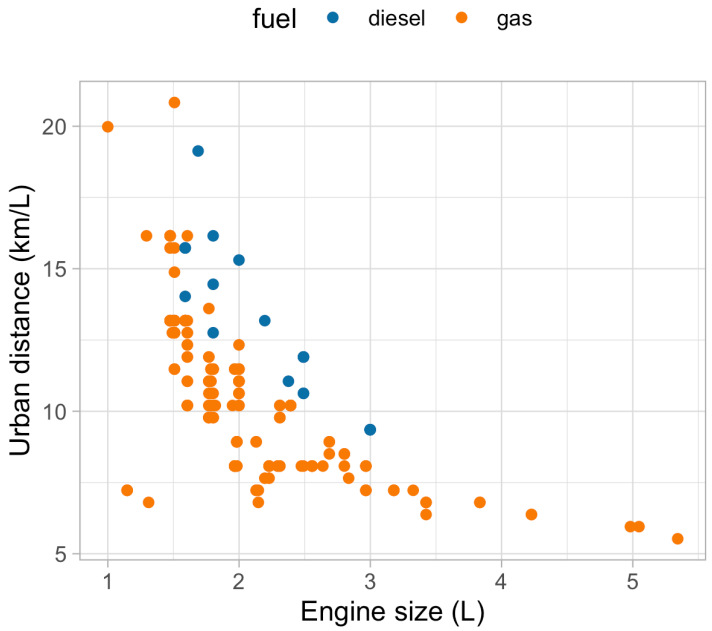

- Let us consider the variables `city.distance` (y), `engine.size` (x) and `fuel` (z).
- A **simple linear regression**

$$Y_i = \beta_1 + \beta_2 x_i + \epsilon_i, \qquad i = 1, \dots, n,$$

  could be easily fit by least squares...
- ... but the plot suggests that the relationship between `city.distance` and `engine.size` is **not** well approximated by a **linear** function.
- ... and also that `fuel` has a non-negligible effect on the response.

Regression models

Linear models

- Let us consider again the variables `city.distance` (y), `engine.size` (x) and `fuel` (z).
- Which function $f(x, z; \beta)$ should we choose?

Matrix notation

- The **response random variables** are collected in the random vector $\boldsymbol{Y} = (Y_1, \ldots, Y_n)^T$, whose **observed realization** is $\boldsymbol{y} = (y_1, \ldots, y_n)^T$.
- The **design matrix** is a $n \times p$ matrix, comprising the covariate's values, defined by

$$\boldsymbol{X} = \begin{bmatrix} x_{11} & \cdots & x_{1p} \\ \vdots & \ddots & \vdots \\ x_{n1} & \cdots & x_{np} \end{bmatrix}.$$

Linear regression: estimation I

- The optimal set of coefficients $\hat{\beta}$ is the minimizer of the **least squared criterion**

$$D(\beta) = (\boldsymbol{y} - \boldsymbol{X}\beta)^T(\boldsymbol{y} - \boldsymbol{X}\beta) = ||\boldsymbol{y} - \boldsymbol{X}\beta||^2,$$

  also known as **residual sum of squares (RSS)**, where

$$||\boldsymbol{y}|| = \sqrt{y_1^2 + \cdots + y_n^2},$$

  denotes the **Euclidean norm**.

Linear regression: estimation II

- In matrix notation, the predicted values can be obtained as

$$\hat{\boldsymbol{y}} = \boldsymbol{X}\hat{\beta} = \boldsymbol{H}\boldsymbol{y}, \qquad \boldsymbol{H} = \boldsymbol{X}(\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{X}^T,$$

  where $\boldsymbol{H}$ is a $n \times n$ **projection matrix** matrix sometimes called **hat matrix**. The matrix is idempotent, meaning that $\boldsymbol{H} = \boldsymbol{H}^T$ and $\boldsymbol{H}^2 = \boldsymbol{H}$.

Linear regression: inference

- Recall that the errors $\boldsymbol{\epsilon}$ have zero mean $\mathbb{E}(\boldsymbol{\epsilon}) = 0$ and are **uncorrelated** $\text{var}(\boldsymbol{\epsilon}) = \sigma^2 I_n$.
- Then, the estimator $\hat{\beta}$ is **unbiased** $\mathbb{E}(\hat{\beta}) = \beta$ and its **variance** is $\text{var}(\hat{\beta}) = \sigma^2(\boldsymbol{X}^T\boldsymbol{X})^{-1}$. Since σ^2 is also unknown, we can estimate the variances of $\hat{\beta}$ as follows:

$$\widehat{\text{var}}(\hat{\beta}) = \boldsymbol{s}^2(\boldsymbol{X}^T\boldsymbol{X})^{-1}.$$

- The **standard errors** of the components of $\hat{\beta}$ correspond to the square root of the diagonal of the above covariance matrix.

Linear regression: diagnostic

- The diagonal elements $h_i \in [0, 1]$ of the matrix $\boldsymbol{H}$ are called **leverages** and it holds

$$\mathrm{var}(\hat{Y}_i) = \sigma^2 h_i, \qquad \mathrm{var}(Y_i - \hat{Y}_i) = \sigma^2(1 - h_i), \qquad \mathrm{cor}(Y_i, \hat{Y}_i) = \sqrt{h_i}.$$

  The leverage h_i determines the **precision** with which $\hat{Y}_i$ predicts Y_i. For large h_i close to 1, $\mathrm{cor}(Y_i, \hat{Y}_i) \approx 1$, therefore changes of a single point Y_i leads to significant changes in $\hat{Y}_i$.

- Leverages also appear in the definition of **standardized residuals**:

$$\tilde{r}_i = \frac{r_i}{\sqrt{s^2(1 - h_i)}} = \frac{y_i - \boldsymbol{x}_i^T \hat{\beta}}{\sqrt{s^2(1 - h_i)}},$$

  where $r_i = y_i - \boldsymbol{x}_i^T \hat{\beta}$ are the (raw) **residuals**.

Leverages, outliers and influence points

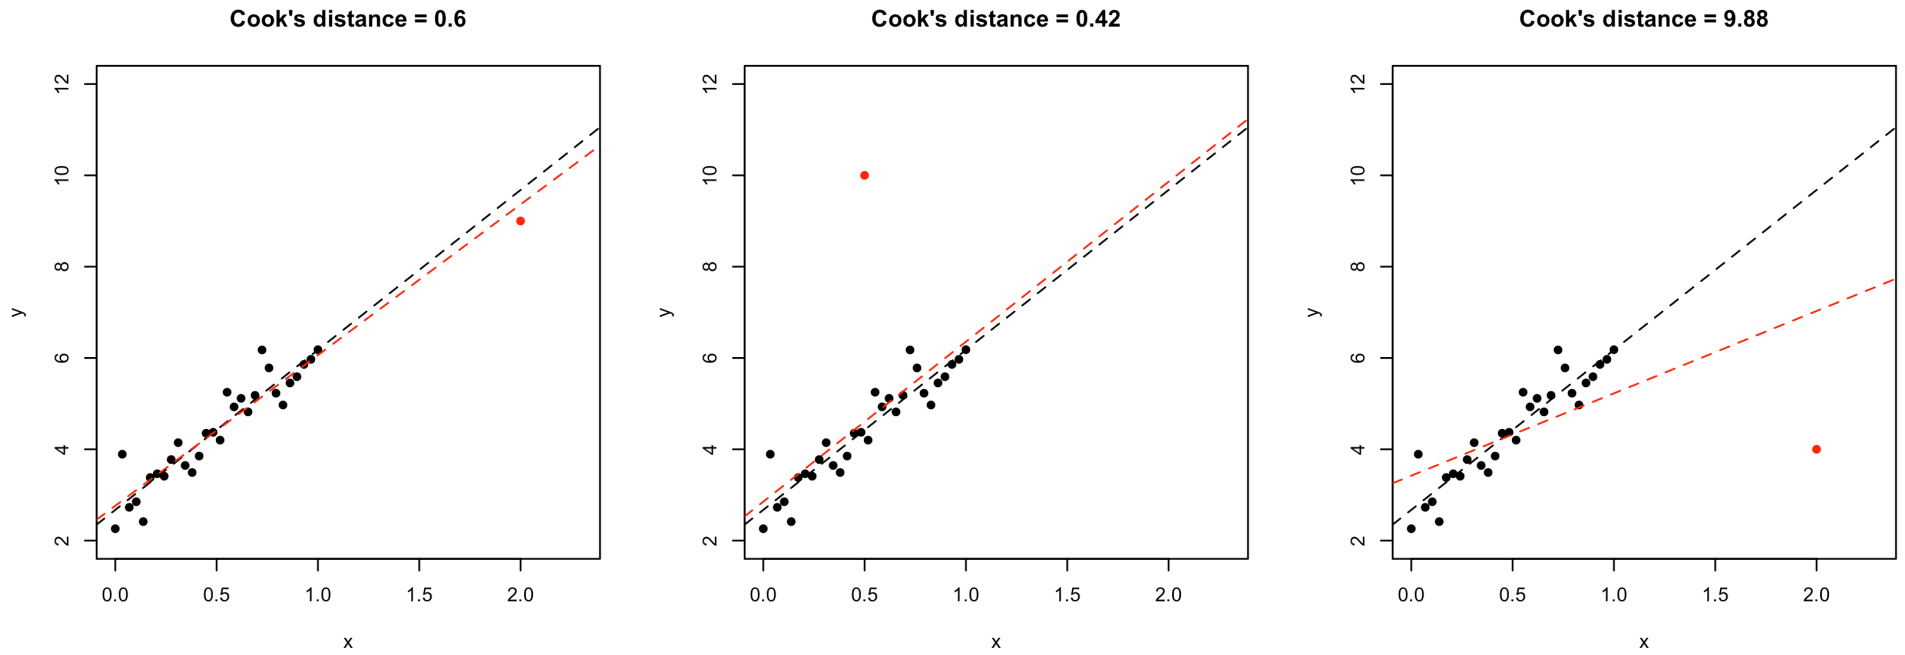

- **Left plot**: leverage, not outlier. **Central plot**: outlier, not leverage. **Right plot**: influence point = leverage + outlier.

A first model: estimated coefficients

- Our first attempt for predicting `city.distance` (y) via `engine.size` (x) and `fuel` (z) is:

$$f(x, z; \beta) = \beta_1 + \beta_2 x + \beta_3 x^2 + \beta_4 x^3 + \beta_5 I(z = \textbf{gas}).$$

- We obtain the following **summary** for the regression coefficients $\hat{\beta}$.

term	estimate	std.error	statistic	p.value
(Intercept)	28.045	3.076	9.119	0.000
engine.size	-10.980	3.531	-3.109	0.002
engine.size^2	2.098	1.271	1.651	0.100
engine.size^3	-0.131	0.139	-0.939	0.349
fuel_gas	-3.214	0.427	-7.523	0.000

- Moreover, the coefficient R^2 and the residual standard deviation s are:

r.squared	sigma	deviance
0.5973454	1.790362	634.6687

Home page

A first model: fitted values

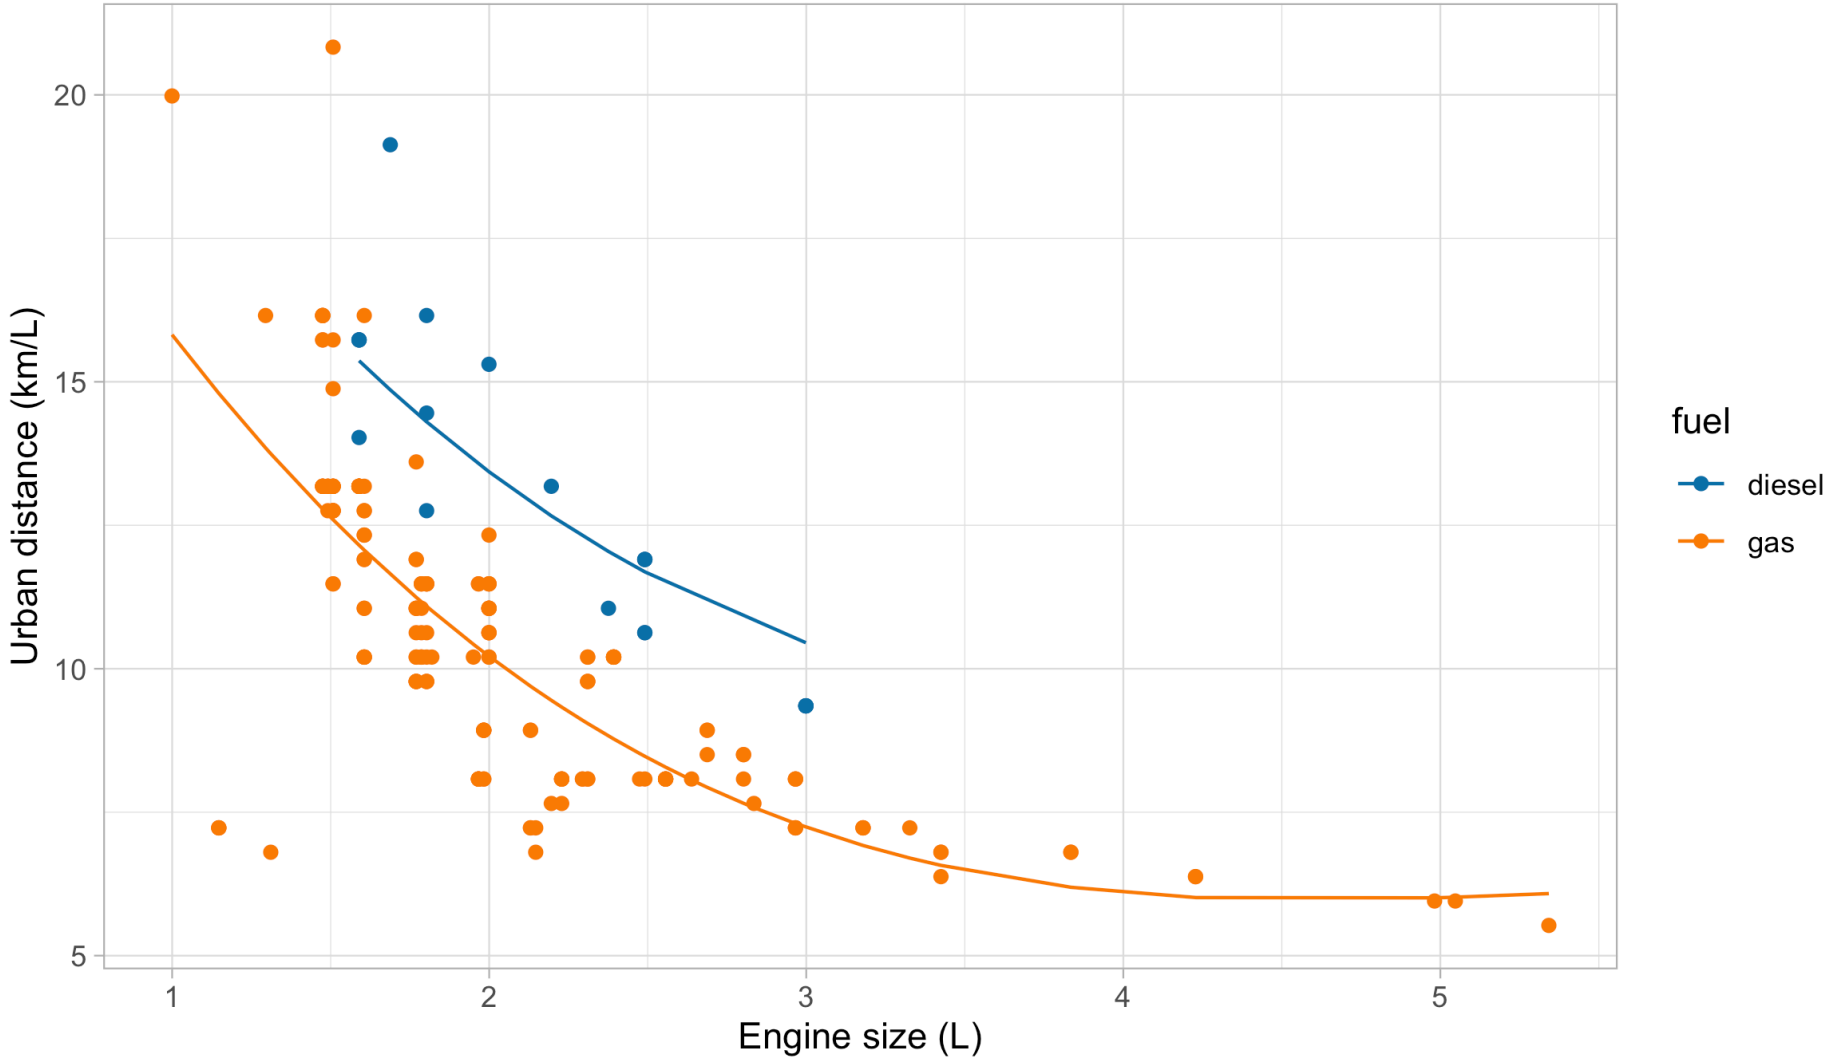

Home page

A first model: graphical diagnostics

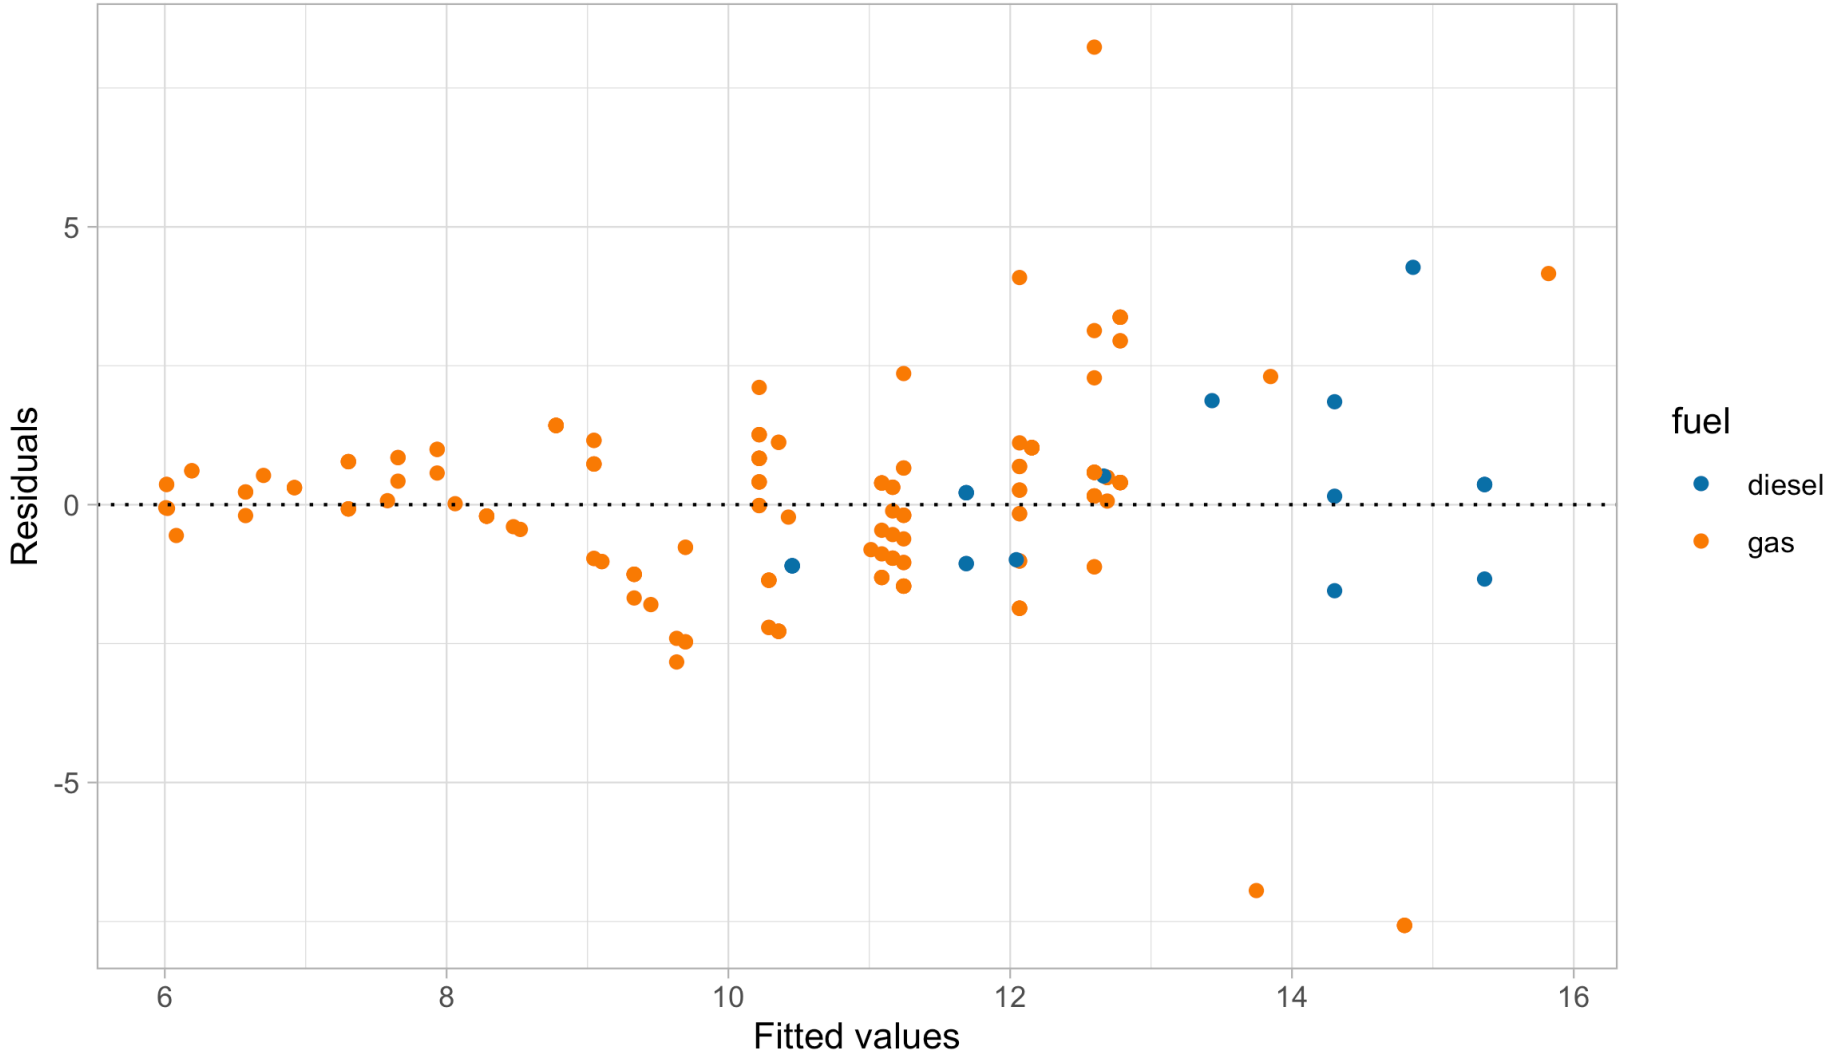

Home page

Comments and criticisms

- Is this a good model?
- The overall fit **seems satisfactory** at first glance, especially if we aim at predicting the urban distance of cars when average engine size (i.e., between $1.5L$ and $3L$).

Linear models and non-linear patterns

- A significant advantage of linear models is that they can describe non-linear relationships via **variable transformations** such as polynomials, logarithms, etc.

Second model: fitted values

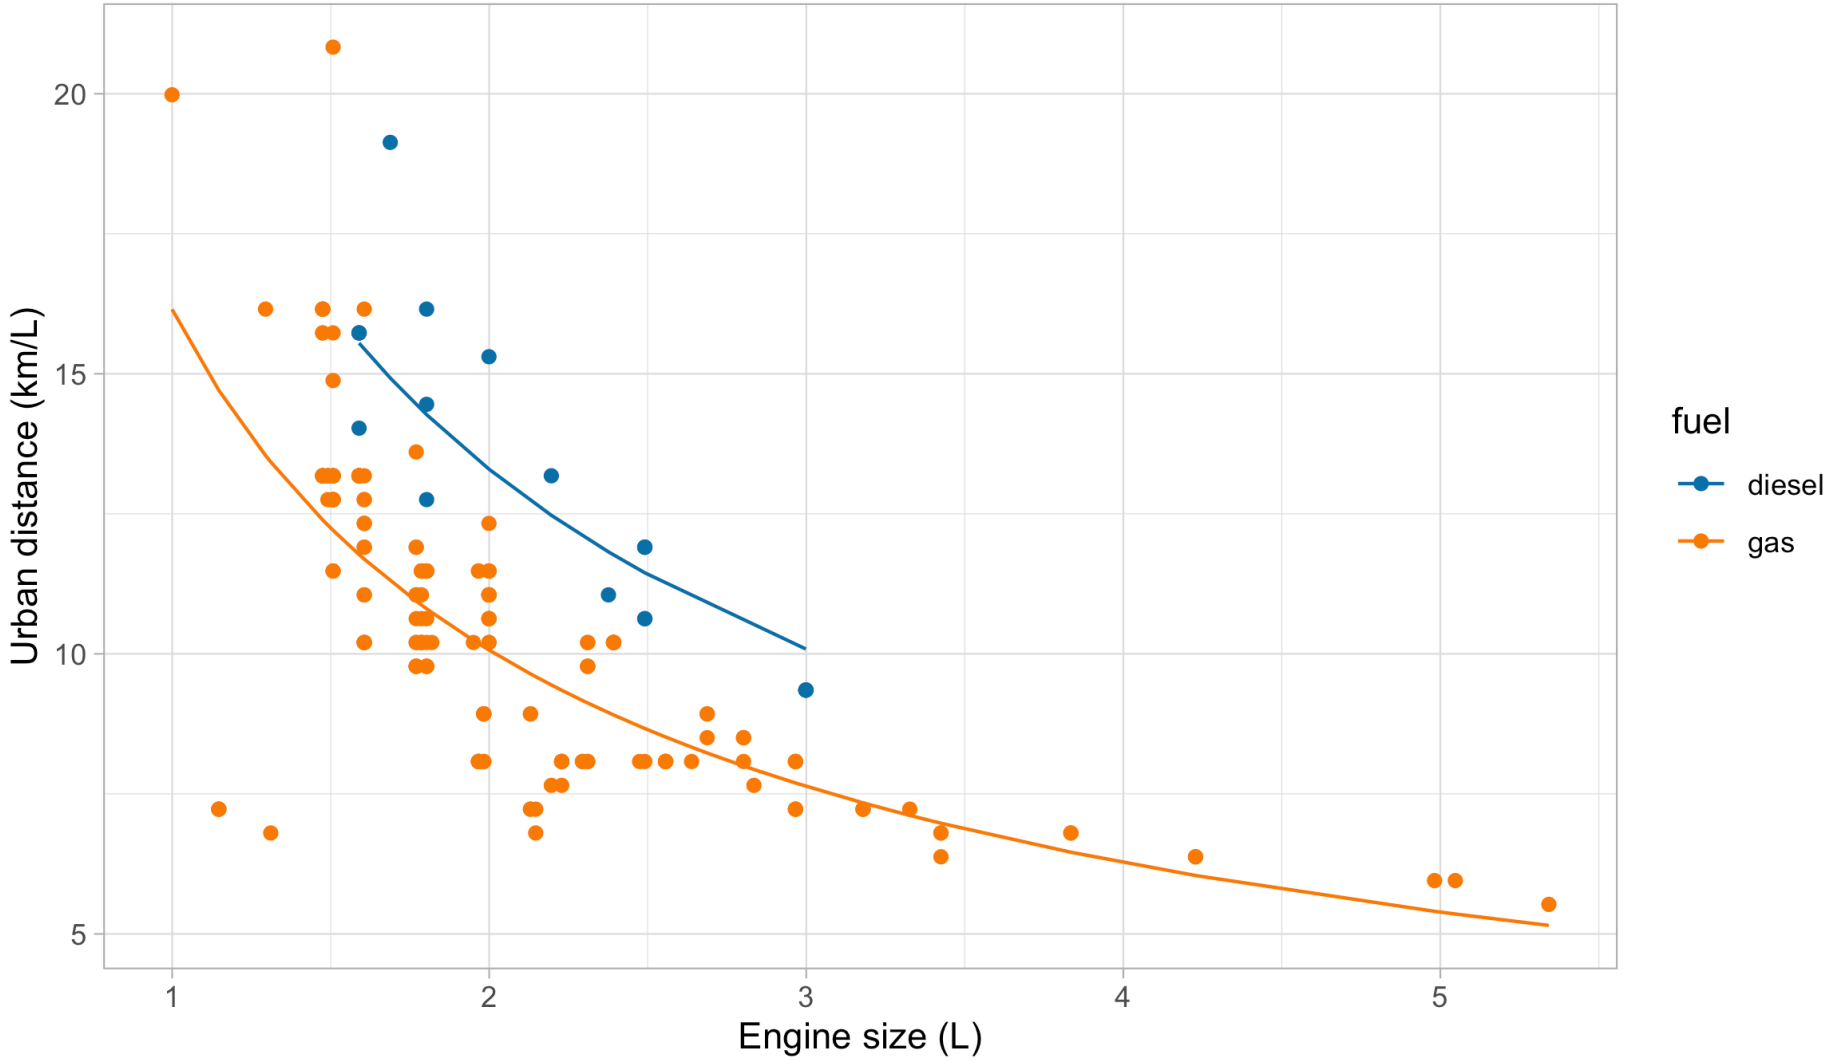

Home page

Second model: graphical diagnostics

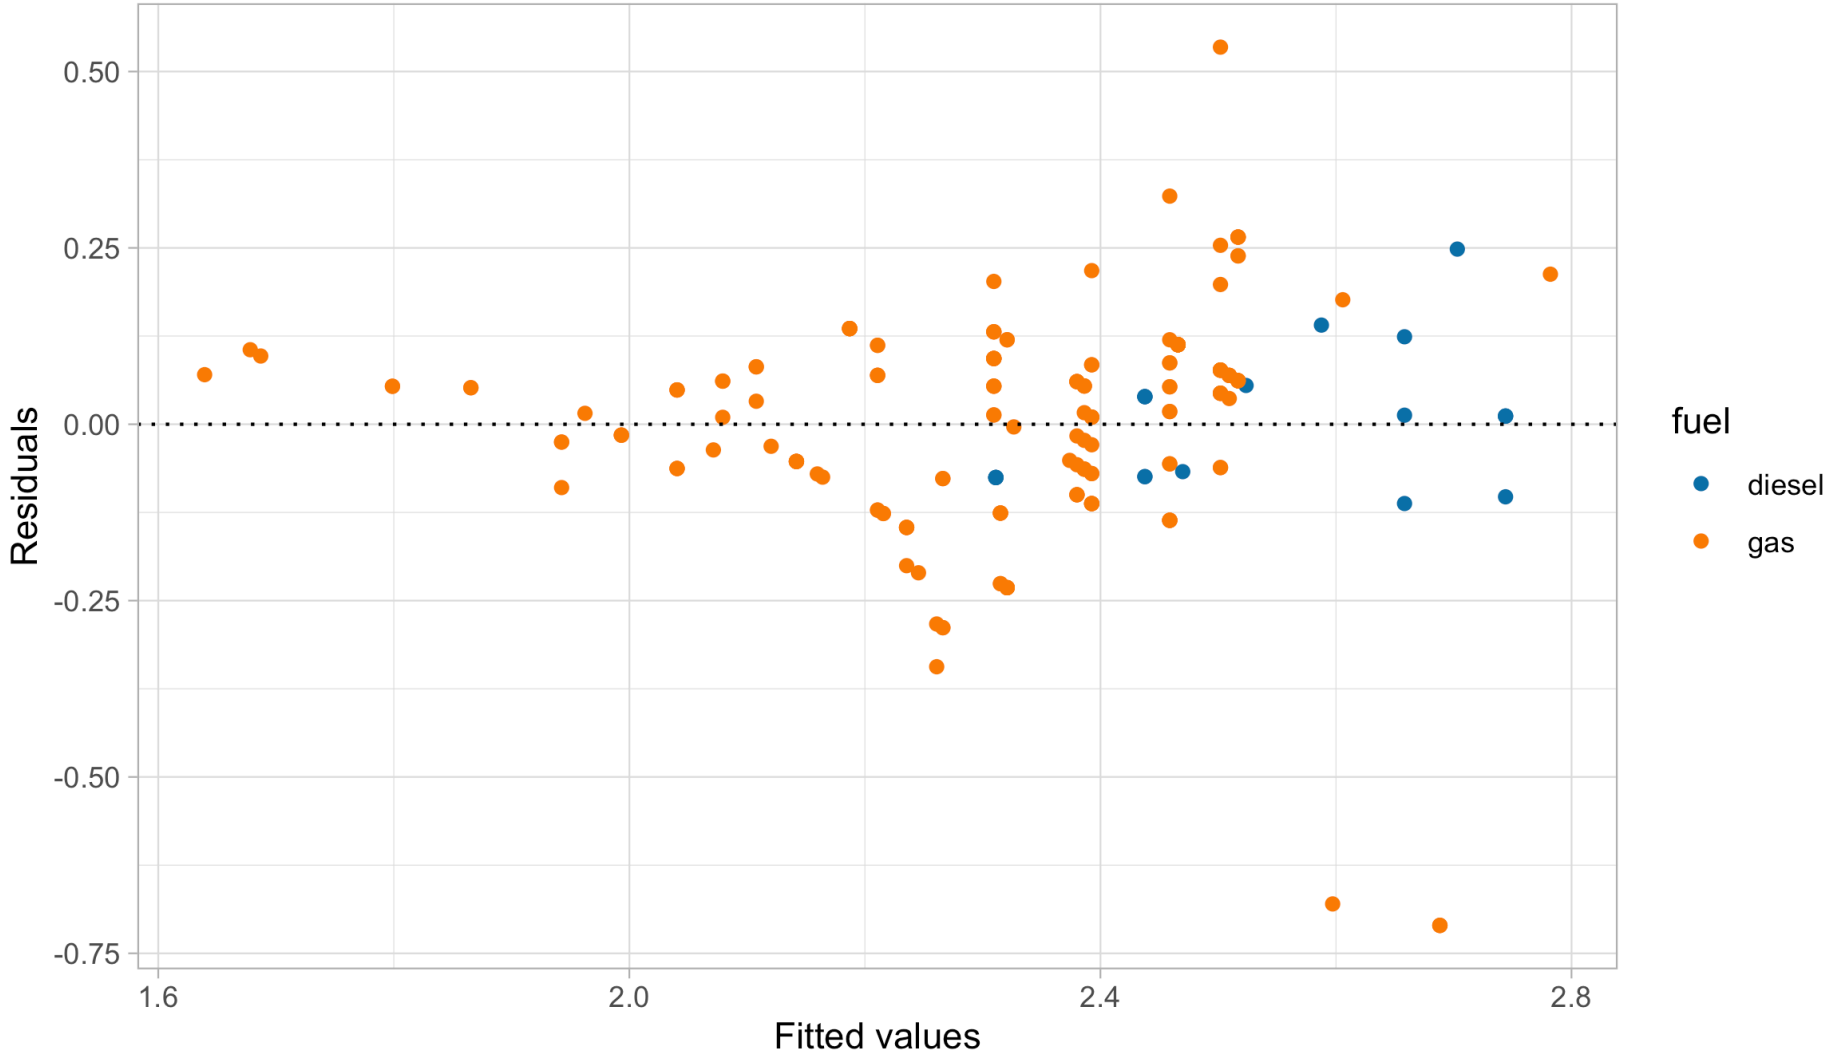

Home page

Comments and criticisms

- The **goodness of fit** indices are the following:

r.squared.original	r.squared	sigma	deviance
0.5847555	0.6196093	0.1600278	5.121777

- Do not mix **apple** and **oranges**! Compare R^2s only if they refer to the same scale!

A third model: additional variables

- Let us consider **two additional variables**: `curb.weight` (w) and `n.cylinders` (v).
- A richer model, therefore, could be:

$$\log Y_i = \beta_1 + \beta_2 \log x_i + \beta_3 \log w_i + \beta_4 I(z_i = \texttt{gas}) + \beta_5 I(v_i = 2) + \epsilon_i,$$

for $i = 1, \ldots, n$. The estimates are:

A third model: graphical diagnostics

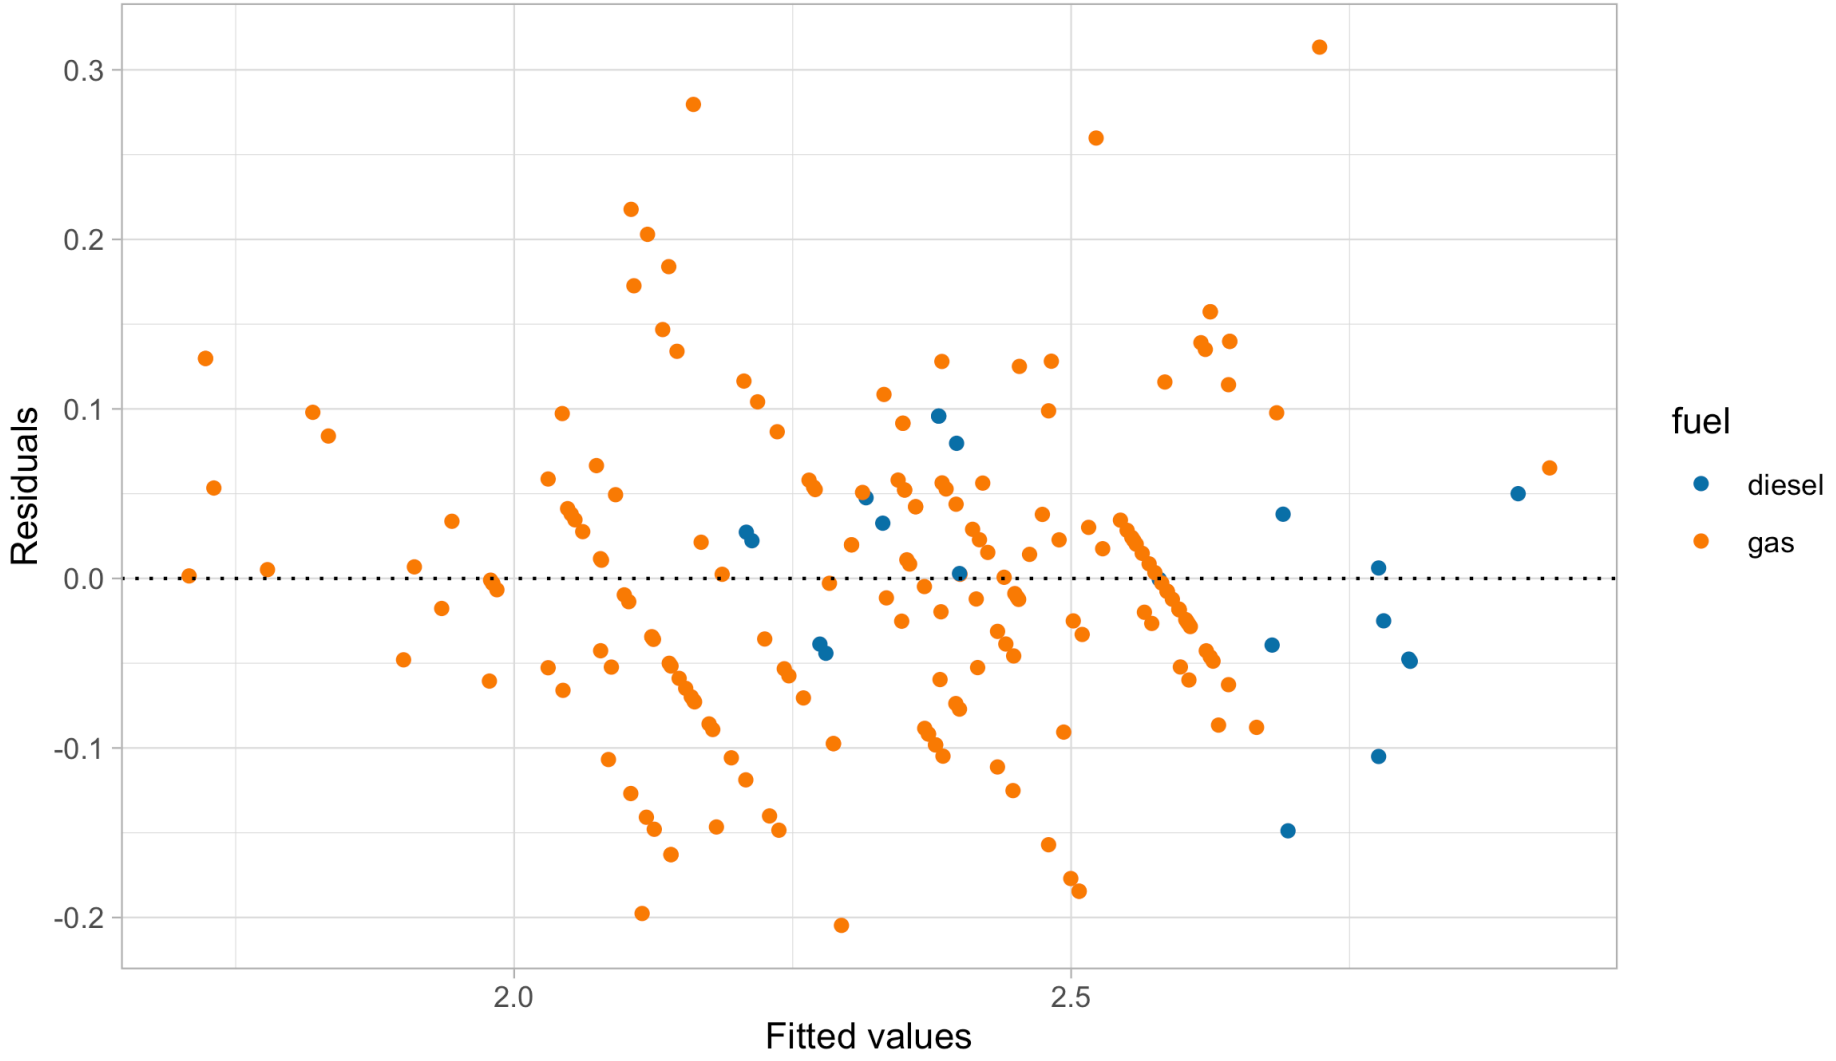

Home page

Comments and criticisms

- The goodness of fit greatly **improved**:

r.squared.original	r.squared	sigma	deviance
0.869048	0.8819199	0.0896089	1.589891

- In this third model, we handled the **outliers** appearing in the residual plots, which it turns out are identified by the group of cars having 2 cylinders.
- The diagnostic plots are also very much improved, although still not perfect.
- The estimates are coherent with our expectations, based on common knowledge. Have a look at the book (Azzalini and Scarpa (2012)) for a detailed explanation of β_4!
- The car dataset is available from the textbook (A&S) website:
  - Dataset http://azzalini.stat.unipd.it/Book-DM/auto.dat
  - Variable description http://azzalini.stat.unipd.it/Book-DM/auto.names

Misspecification and remedies

Assumptions and misspecification

Classical assumptions of linear models

- **(A.1)** **Linear structure**, namely $\boldsymbol{Y} = \boldsymbol{X}\beta + \boldsymbol{\epsilon}$ with $\mathbb{E}(\boldsymbol{\epsilon}) = 0$, implying $\mathbb{E}(\boldsymbol{Y}) = \boldsymbol{X}\beta$. [1]
- **(A.2)** **Homoschedasticity** and **uncorrelation** of the errors, namely $\mathrm{var}(\boldsymbol{\epsilon}) = \sigma^2 I_n$.
- **(A.3)** **Gaussianity**, namely $\boldsymbol{\epsilon} \sim \mathrm{N}_n(0, \sigma^2 I_n)$. In other words, the errors $\epsilon_i \overset{\text{iid}}{\sim} N(0, \sigma^2)$ are iid Gaussian random variables with zero mean and variance σ^2.

It is also commonly asked that $\mathrm{rk}(\boldsymbol{X}) = p$, otherwise the model is not identifiable.

- If one of the above assumptions is violated, it is not necessarily a huge problem, because
  - the OLS estimator $\hat{\beta}$ is fairly **robust** to misspecification;
  - simple **fixes** (variable transformations, standard error corrections) are available.

1. If the intercept is included in $\boldsymbol{X}$, the errors automatically satisfy the property $\mathbb{E}(\boldsymbol{\epsilon}) = 0$.

Robust estimation and assumptions

- A plane can still fly with one of its **engines on fire**, but this is hardly an appealing situation.
- Similarly, robust estimators may work under **model misspecification**, but this does not mean we should neglect **checking** whether the original **assumptions** hold.

Non-normality of the errors I 📖

- Let us consider the case in which assumptions **(A.1)**-**(A.2)** are **valid** but **(A.3)** **is not**, that is $\mathbb{E}(\boldsymbol{\epsilon}) = 0$ and $\mathrm{var}(\boldsymbol{\epsilon}) = \sigma^2 I_n$, but ϵ does **not** follow **a Gaussian** distribution.
- For example, ϵ_i may follow a Laplace distribution, a skew-Normal, a logistic distribution, a Student's t distribution, etc.

Non-normality of the errors II

- When the errors are non Gaussian the **exact inferential results** are not valid. In particular $\hat{\beta}$ does not follow anymore a Gaussian distribution.
- However, a **central limit theorem** can be invoked under very mild conditions on the design matrix $\boldsymbol{X}$.
- Thus, when the sample size n is large enough, then the following **approximation** holds

$$\hat{\beta} \dot{\sim} \mathrm{N}_p(\beta, \sigma^2(\boldsymbol{X}^T\boldsymbol{X})^{-1}),$$

from which **confidence intervals** and **test statistics** can be obtained as usual. The approximation is **excellent** if the errors are **symmetric** around 0.

Heteroschedasticity of the errors I 📖

- Suppose now that the linearity assumption **(A.1)** is valid but **homoschedasticity** of the errors **(A.2)** is **not**. Instead, we consider **heteroschedastic errors**:

$$\text{var}(\boldsymbol{\epsilon}) = \boldsymbol{\Sigma}, \quad \text{or equivalenty} \quad \text{var}(Y_i) = \sigma_i^2, \quad i = 1, \dots, n$$

  where $\boldsymbol{\Sigma} = \text{diag}(\sigma_1^2, \dots, \sigma_n^2)$ is a diagonal matrix with positive entries.

- The OLS estimator is still **unbiased**, with a **modified covariance** structure[1]

$$\mathbb{E}(\hat{\beta}) = \beta, \qquad \text{var}(\hat{\beta}) = (\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{X}^T\boldsymbol{\Sigma}\boldsymbol{X}(\boldsymbol{X}^T\boldsymbol{X})^{-1}.$$

  If in addition we assume Gaussianity of the errors, that is $\boldsymbol{\epsilon} \sim \text{N}_n(0, \boldsymbol{\Sigma})$, then

$$\hat{\beta} \sim \text{N}_p(\beta, (\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{X}^T\boldsymbol{\Sigma}\boldsymbol{X}(\boldsymbol{X}^T\boldsymbol{X})^{-1}).$$

  Under suitable but mild conditions on $\boldsymbol{X}$ and $\boldsymbol{\Sigma}$, the estimator is also **consistent**.

1. These results are valid even when the matrix $\boldsymbol{\Sigma}$ is non-diagonal. This is useful to model correlated responses.

Heteroschedasticity of the errors II

The OLS estimator in presence of heteroschedasticity still gives a **good point estimate**. However, the OLS estimator is **not efficient** and the classical **standard errors** are **wrong**.

- A potential approach is to **accept the inefficiency** of the OLS estimator in this scenario and **correct** the standard errors.
- The elements of $\boldsymbol{\Sigma}$ are **unknown**, but we can estimate them from the data. Note that

$$\mathrm{var}(r_i) = \mathrm{var}(y_i - \boldsymbol{x}_i^T \hat{\beta}) = \sigma_i^2 (1 - h_i),$$

  suggesting the **estimate** $\hat{\sigma}_i^2 = r_i^2/(1 - h_i)$.
- This leads to the so-called **sandwich estimator** of the covariance matrix:

$$\widehat{\mathrm{var}}(\hat{\beta}) = (\boldsymbol{X}^T \boldsymbol{X})^{-1} \boldsymbol{X}^T \hat{\boldsymbol{\Sigma}} \boldsymbol{X} (\boldsymbol{X}^T \boldsymbol{X})^{-1},$$

  where $\hat{\boldsymbol{\Sigma}} = \mathrm{diag}(\hat{w}_1, \ldots, \hat{w}_n)$ and $\hat{w}_i = r_i^2/(1 - h_i)$.
- These are known as **White's** heteroscedasticity-consistent **standard errors**. [1]

1. White originally proposed the simpler version $\hat{\sigma}_i^2 = r_i^2$. Another variant is $\hat{\sigma}_i^2 = r_i^2/(1 - h_i)^2$.

Weighted least squares I 📖

- Let us consider again the case of **heteroschedastic errors**:

$$\text{var}(\boldsymbol{\epsilon}) = \sigma^2 \boldsymbol{\Omega}^{-1}, \quad \text{or equivalenty} \quad \text{var}(Y_i) = \sigma_i^2 = \frac{\sigma^2}{\omega_i}, \quad i = 1, \dots, n$$

  where $\boldsymbol{\Omega} = \text{diag}(\omega_1, \dots, \omega_n)$ are positive **weights**. However, here we assume that the weights $\omega_1, \dots, \omega_n$ are **known**, a common situation in survey design.

- Let us define the **standardized** quantities:

$$\boldsymbol{Y}^* = \boldsymbol{\Omega}^{1/2} \boldsymbol{Y}, \qquad \boldsymbol{X}^* = \boldsymbol{\Omega}^{1/2} \boldsymbol{X}.$$

  This is equivalent to say that $Y_i^* = \sqrt{\omega_i} Y_i$ and $x_{ij}^* = \sqrt{\omega_i} x_{ij}$. Then, it is easy to show that

$$\mathbb{E}(\boldsymbol{Y}^*) = \boldsymbol{X}^* \beta, \qquad \text{var}(\boldsymbol{Y}^*) = \sigma^2 \boldsymbol{\Omega}^{1/2} \boldsymbol{\Omega}^{-1} \boldsymbol{\Omega}^{1/2} = \sigma^2 I_n,$$

  namely the **assumptions (A.1)** and **(A.2)** are valid in the **transformed scale**.

- In other words, **after** a suitable **transformation**, we reconducted the problem to a **standard linear model**.

Weighted least squares II 📖

- Thus an estimator for β, based on the transformed data, is obtained minimizing the deviance

$$D_{\text{wls}}(\beta) = (\boldsymbol{y}^* - \boldsymbol{X}^*\beta)^T(\boldsymbol{y}^* - \boldsymbol{X}^*\beta) = (\boldsymbol{y} - \boldsymbol{X}\beta)^T\boldsymbol{\Omega}(\boldsymbol{y} - \boldsymbol{X}\beta)$$
$$= \sum_{i=1}^{n} \omega_i(y_i - \boldsymbol{x}_i^T\beta)^2.$$

  which is a **weighted** version of the original **quadratic loss**, with **high weight = low variance**.

- The resulting OLS estimate minimizing $D_{\text{wls}}(\beta)$ in the transformed and original scales is

$$\hat{\beta}_{\text{wls}} = [(\boldsymbol{X}^*)^T\boldsymbol{X}^*]^{-1}(\boldsymbol{X}^*)^T\boldsymbol{y}^* = (\boldsymbol{X}^T\boldsymbol{\Omega}\boldsymbol{X})^{-1}\boldsymbol{X}^T\boldsymbol{\Omega}\boldsymbol{y}$$

  and it is referred to as **weighted least squares** estimator of β.

- Such an estimator is **unbiased** and **efficient** (**BLUE**), with

$$\mathbb{E}(\hat{\beta}_{\text{wls}}) = \beta, \qquad \text{var}(\hat{\beta}_{\text{wls}}) = \sigma^2(\boldsymbol{X}^T\boldsymbol{\Omega}\boldsymbol{X})^{-1}.$$

  Moreover, if $\boldsymbol{\epsilon} \sim \text{N}_n(0, \sigma^2\boldsymbol{\Omega}^{-1})$ it also coincides with the **maximum likelihood** estimator.

Variable transformations

- Another remedy for **misspecification** was already applied in the analysis of the car dataset, namely through **variable transformation**.

While the model may have been incorrectly specified for the original data, it could become **appropriate** once the **transformations** are considered, namely

$$g(Y_i) = h_1(\boldsymbol{x}_i)\beta_1 + \cdots + h_p(\boldsymbol{x}_i)\beta_p + \epsilon_i, \qquad i = 1, \ldots, n,$$

where $g(\cdot)$ and $h_j(\cdot)$ for $j = 1, \ldots, p$ are **non-linear** and **known** functions.

Box-Cox transform

Box-Cox transform

If the data are y_i are **positive**, we may consider a **parametric class** of transformations:

$$g_\lambda(y) = \frac{y^\lambda - 1}{\lambda}, \qquad \lambda \neq 0.$$

and $g_\lambda(y) = \log y$ when $\lambda = 0$. This is the celebrated **Box-Cox transform**.

The case $\lambda = 1$ corresponds to no transformation, $\lambda = 1/2$ to the square root, $\lambda = 0$ to the logarithm, and $\lambda = -1$ to the reciprocal.

Box-Cox transform: derivation I 📖

- By assumption, the distribution of the **transformed data** $\boldsymbol{Z}_\lambda = (g_\lambda(Y_1), \ldots, g_\lambda(Y_n))^T$ is Gaussian, therefore their joint density is

$$f_Z(\boldsymbol{z}_\lambda) = \frac{1}{(2\pi\sigma^2)^{n/2}} \exp\left\{-\frac{1}{2\sigma^2}(\boldsymbol{z}_\lambda - \boldsymbol{X}\beta)^T(\boldsymbol{z}_\lambda - \boldsymbol{X}\beta)\right\}.$$

- Using standard tools of probability theory, we can obtain the density of the **original data**:

$$f_Y(\boldsymbol{y}) = f_Z(g_\lambda(y_1), \ldots, g_\lambda(y_n)) \prod_{i=1}^n \left|\frac{\partial g_\lambda(y_i)}{\partial y_i}\right|, \quad \text{where} \quad \left|\frac{\partial g_\lambda(y_i)}{\partial y_i}\right| = y_i^{\lambda-1}.$$

  The additional term is the determinant of the **Jacobian** of the transformation.

- The **log-likelihood** therefore is

$$\ell(\beta, \sigma^2, \lambda) = -\frac{n}{2}\log\sigma^2 - \frac{1}{2\sigma^2}(\boldsymbol{z}_\lambda - \boldsymbol{X}\beta)^T(\boldsymbol{z}_\lambda - \boldsymbol{X}\beta) + (\lambda - 1)\sum_{i=1}^n \log y_i.$$

Box-Cox transform: derivation II 📖

- Note that, for any given value of λ, the maximum likelihood estimates are

$$\hat{\beta}_\lambda = (\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{X}^T\boldsymbol{z}_\lambda, \qquad \hat{\sigma}^2_\lambda = \frac{1}{n}(\boldsymbol{z}_\lambda - \boldsymbol{X}\hat{\beta}_\lambda)^T(\boldsymbol{z}_\lambda - \boldsymbol{X}\hat{\beta}_\lambda),$$

- We can **plug-in** the above estimates into the log-likelihood. This gives the **profile log-likelihood** for λ, which admits a very simple expression:

$$\ell_P(\lambda) = \ell(\hat{\beta}_\lambda, \hat{\sigma}^2_\lambda, \lambda) = -\frac{n}{2}\log\hat{\sigma}^2_\lambda + (\lambda - 1)\sum_{i=1}^{n}\log y_i,$$

  which must be **numerically maximized** over λ, e.g. using `optim`.

- The optimal value $\hat{\lambda} = \arg\max \ell_P(\lambda)$, as well as a confidence interval for it, may offer guidance in choosing the right transformation.

Box-Cox transform for the auto dataset

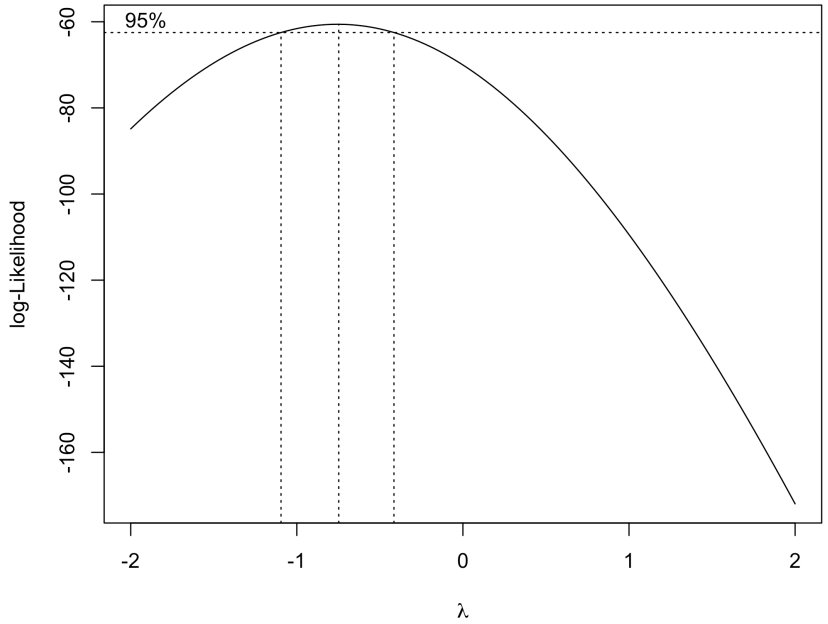

- The **Box-Cox transform** in the auto dataset suggests a **reciprocal** transformation:

$$\frac{1}{Y_i} = \beta_1 + \beta_2 x_i + \beta_3 w_i + \beta_4 I(z_i = \mathbf{gas}) + \beta_5 I(v_i = 2) + \epsilon_i,$$

which is a good alternative to our model based on logarithms of y_i, x_i, and w_i (**but**...).

Home page

A fourth model: graphical diagnostics

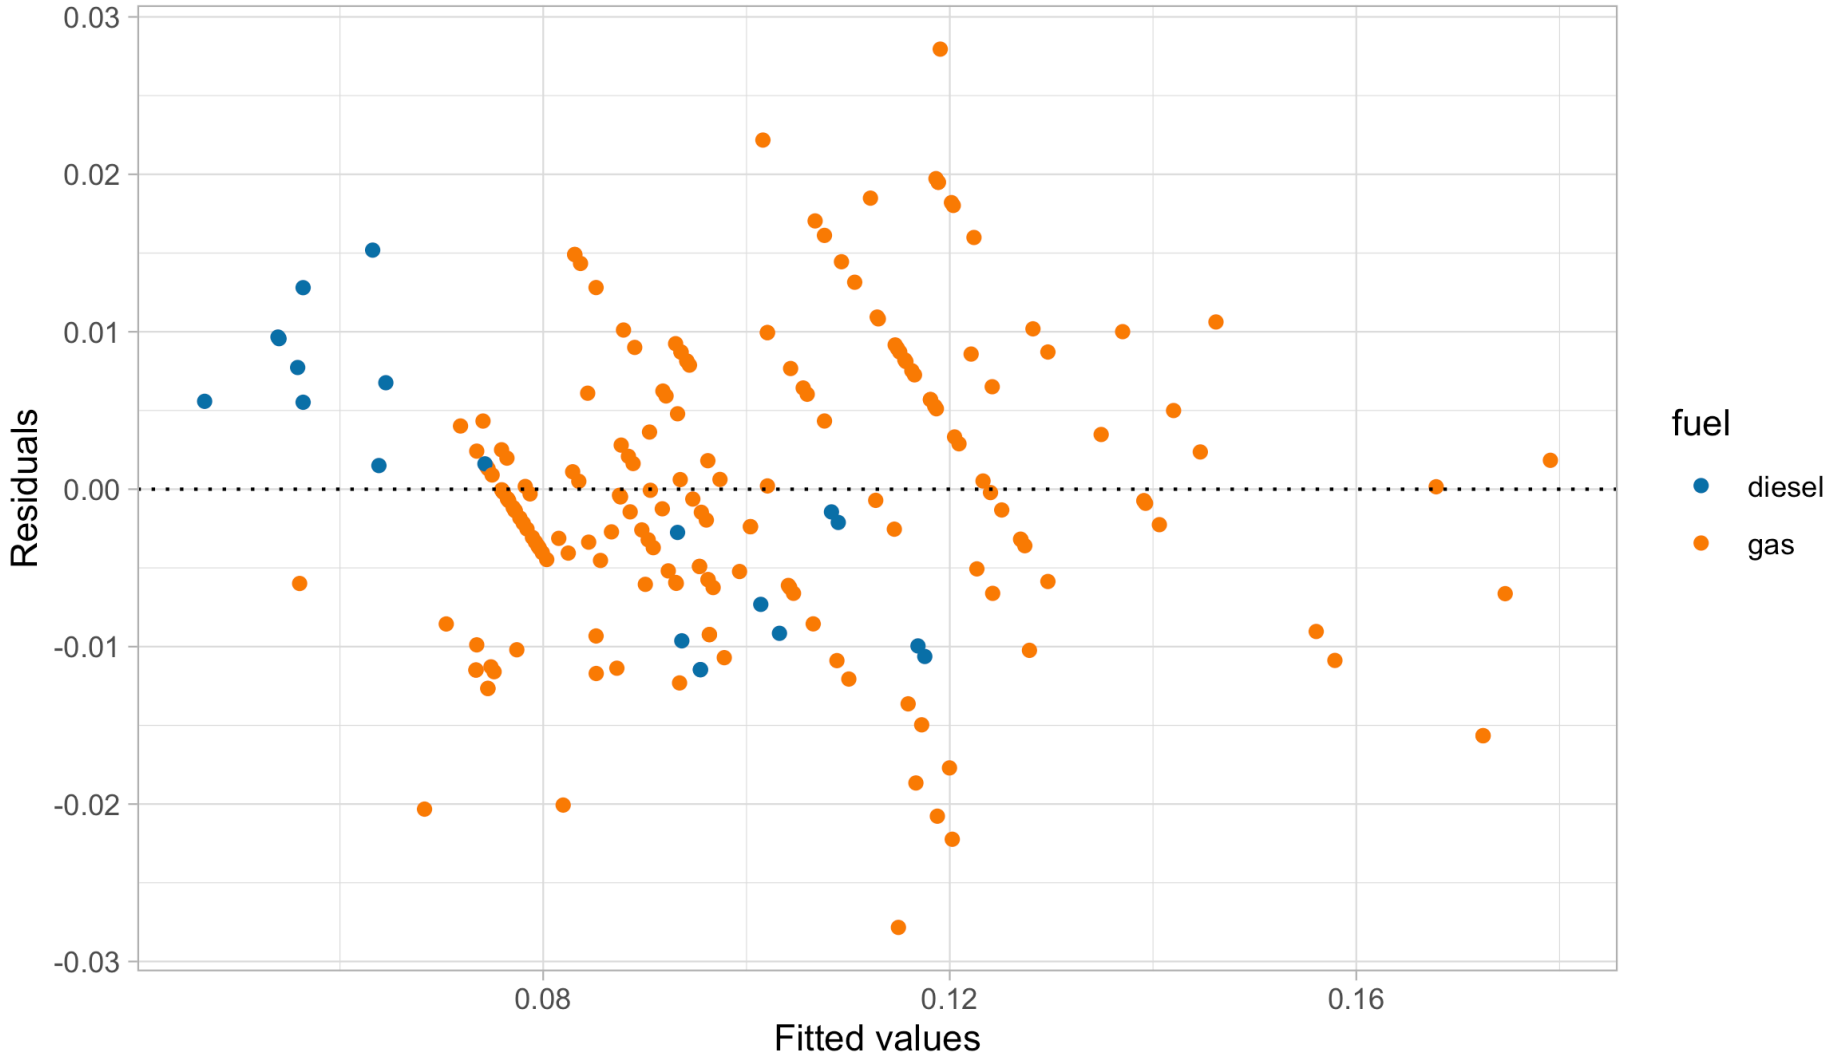

Home page

Variance stabilizing transformations I 📖

- Let $Y_i \sim \text{Poisson}(\mu_i)$ with mean $\mathbb{E}(Y_i) = \mu_i = f(\boldsymbol{x}_i; \beta) = \text{var}(Y_i)$. Note that

$$Y_i \dot{\sim} \text{N}(\mu_i, \mu_i),$$

  is **asymptotically Gaussian** for large values of μ_i. However, data are **heteroschedastic**.

- In modeling count data, we could transform the counts so that, at least **approximately**, the **variance** of $g(Y_i)$ is **constant** and ordinary least squares methods can be used.

Variance stabilizing transformations II 📖

- Let $Y_i \sim \text{Binomial}(\pi_i, m_i)$, with **success probability** $\pi_i = f(\boldsymbol{x}_i; \beta)$ and **trials** m_i. For large values of m_i, the **Gaussian approximation** holds

$$Y_i \dot{\sim} \text{N}(m_i\pi_i, m_i\pi_i(1-\pi_i)).$$

  However, the data are **heteroschedastic**, because $\text{var}(Y_i) = m_i\pi_i(1-\pi_i)$.

- Thus, a **variance stabilizing** transformation in this case is

$$g_{m_i}(y) = \sqrt{m_i}\arcsin\left(\frac{2y}{m_i} - 1\right),$$

  because in fact we have that

$$\text{var}(g_{m_i}(Y_i)) \approx \left(\frac{\sqrt{m_i}}{\sqrt{1-(2\pi_i-1)^2}}\frac{2}{m_i}\right)^2 m_i\pi_i(1-\pi_i) = 1.$$

Limitations of variable transformations I

- Variable transformations are appealing for their simplicity and have a long history in statistics. However, they also have some **drawbacks**.

Limitations of variable transformations II

Suppose $Y_i \sim \text{Binomial}(\pi, m_i)$. The variance stabilizing transformation is not fully satisfactory:

- It **complicates** the **interpretation**, because it models $\mathbb{E}\{g(Y_i)\}$ instead of $\mathbb{E}(Y_i)$;
- It is an **asymptotic approximation**, and is only valid for $m_i \to \infty$.
- The transform depends on m_i, therefore we cannot make predictions for a generic covariate value $\boldsymbol{x}_i$ without knowing the associated m_i.

Besides, this transform is clearly not applicable when $m_i = 1$ and $Y_i \in \{0, 1\}$, a very common problem called **binary regression**.

- If we know that Y_i follows, say, a Bernoulli or a Gamma distribution, then we should use the **appropriate likelihood** rather than a **Gaussian approximation**.
- **Generalized Linear Models** provide a **much more elegant solution** to the above problem.

References

Agresti, A. (2015), *Foundations of Linear and Generalized Linear Models*, Wiley.

Azzalini, A. (2008), *Inferenza statistica*, Springer Verlag.

Azzalini, A., and Scarpa, B. (2012), *Data analysis and data mining: An introduction*, Oxford University Press.

Salvan, A., Sartori, N., and Pace, L. (2020), *Modelli lineari generalizzati*, Springer.